




Introduction to MPLS VPN

Nagendra Kumar, CCIE




Agenda

- **Introduction to MPLS VPN**
 - **MPLS VPN Architecture**
 - **Control Plane Functionality**
 - **Data Plane Functionality**
 - **Basic Configuration**
- **MPLS VPN Troubleshooting Overview**
- **Summary**



MPLS VPNs

Layer 3



What Is a Virtual Private Network?

- VPN is a **set of sites** or **groups** which are allowed to **communicate** with each other
- VPN is defined by **a set of administrative policies**
 - Policies **established** by **VPN customers**
 - Policies could be **implemented** completely by **VPN service providers**
- Flexible inter-site connectivity
 - Ranging from complete to partial mesh
- Sites may be either within the same or in different organizations
 - VPN can be either intranet or extranet
- Site may be in more than one VPN
 - VPNs may overlap
- Not all sites have to be connected to the same service provider
 - VPN can span multiple providers

L2 vs. L3 VPNs

Layer 2 VPNs


- Customer endpoints (CPE) connected via Layer 2 such as Frame Relay DLCI, ATM VC or point-to-point connection
- Provider network is not responsible for distributing site routers as routing relationship is between the customer endpoints
- Provider will need to manually fully mesh end points if any-to-any connectivity is required

Layer 3 VPN


- Customer end points peer with providers' routers @ L3
- Provider network responsible for distributing routing information to VPN sites
- Don't have to manually fully mesh customer endpoints to support any-to-any connectivity



Layer 3 VPNs




IP L3 vs. MPLS L3 VPNs



Overlay VPN

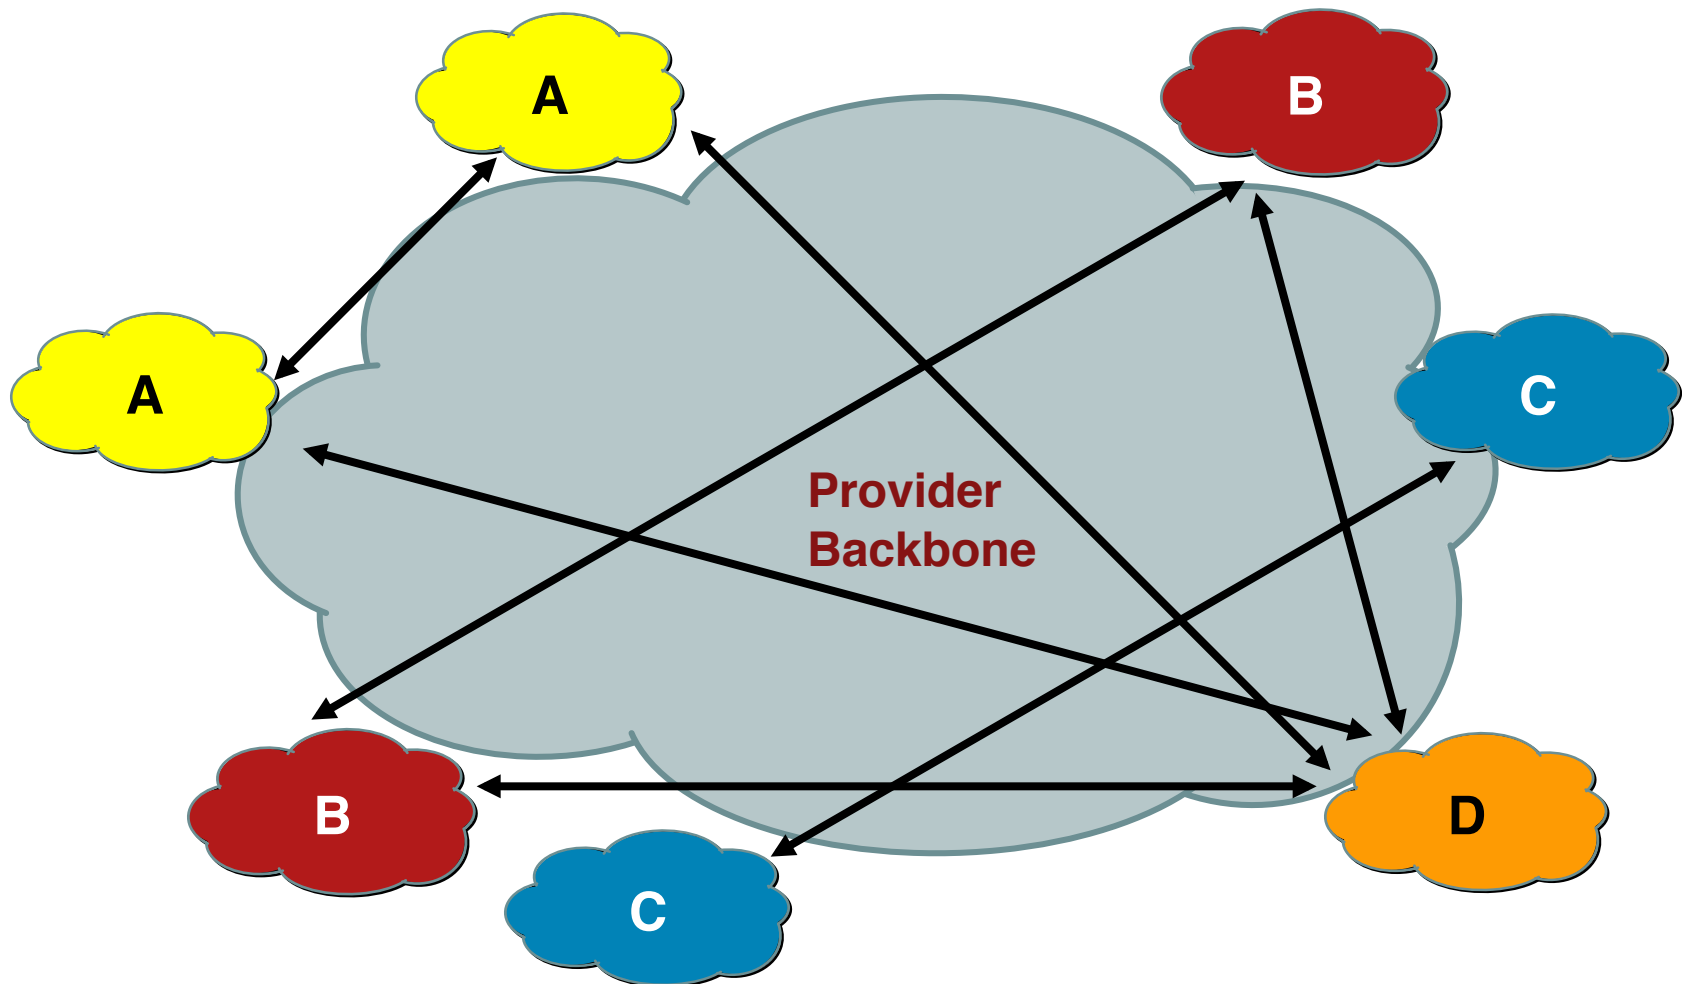
- ACLs, ATM/FR, IP tunnels, IPSec, ...etc. requiring $n*(n-1)$ peering points
- Transport dependent
- Groups endpoints, not groups
- Pushes content outside the network
- Costs scale exponentially
- NAT necessary for overlapping address space
- Limited scaling
- QoS complexity



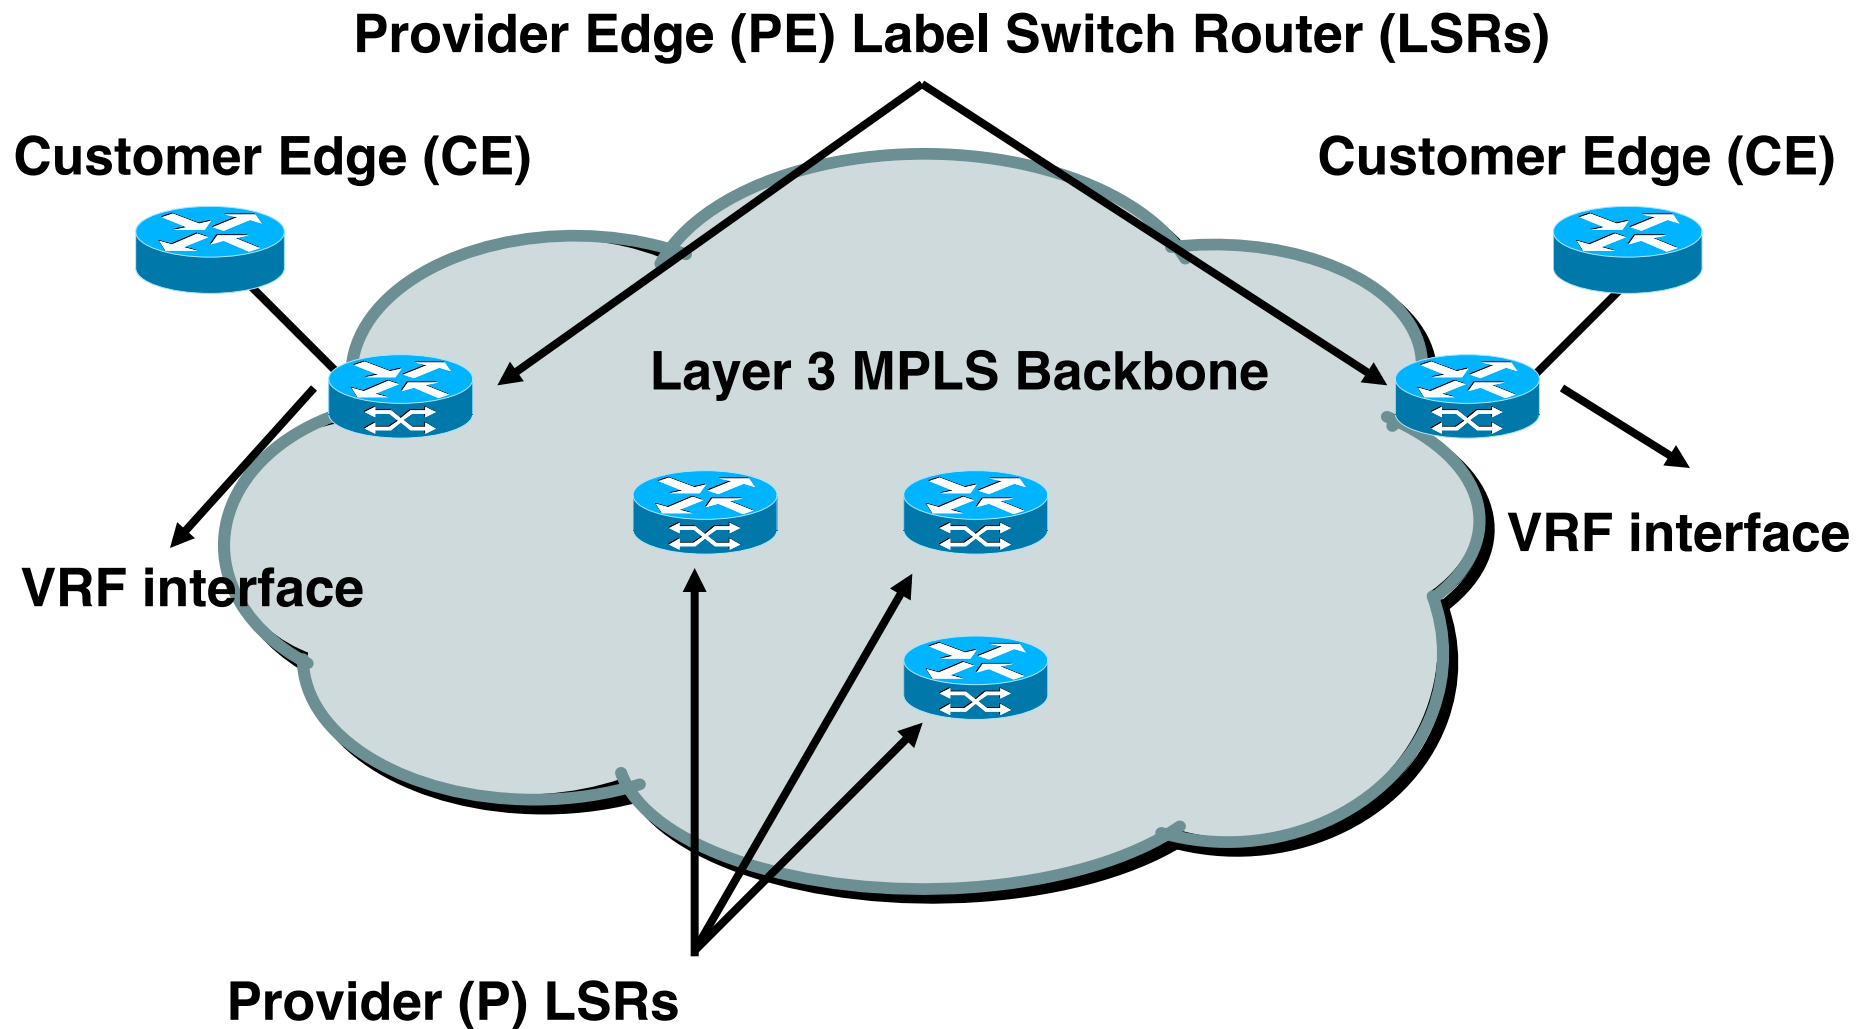
MPLS-Based VPNs

- Point to Cloud single point of connectivity
- Transport independent
- Easy grouping of users and services
- Enables content hosting **inside** the network
- “Flat” cost curve
- Supports private overlapping IP addresses
- Scalable to over millions of VPNs
- Per VPN QoS

VPN Model



MPLS VPN Architecture



MPLS VPN Building Blocks

- MPLS framework (labels) in the core
 - IGP (any)
 - LDP or MPLS Traffic Engineering
- VRF (Virtual Routing/Forwarding) context to keep VPNs separate
 - VRF on PE interface towards CE
 - VRF routing table
 - VRF CEF table

VRF knowledge only needed on edge routers
- RD attached to prefixes to make VPN prefixes unique
 - RD is 64 bits
 - RD allows for overlapping VPN prefixes
- Route targets (ext BGP community) attached to VPN prefixes to allow prefixes to be imported/exported to VPNs
- BGP in the core to advertise VPN prefix and VPN label to all Provider Edge (PE) routers

VRF

- Virtual Routing/Forwarding
- Seperate context for each VPN
 - Seperate RIB per VPN
 - Seperate FIB per VPN
- Each protocol needs to be “VRF-aware” when running across VRF interface
 - e.g. any routing protocol
 - DHCP
 - NAT
 - MIB
 - etc.

VRF Routing Tables

- VRF routing table contains routes that should be available to a particular set of VPN sites
- VRF routing tables support the same set of mechanisms as the standard (default/global) routing table
- There is still the “global” routing table used in the core MPLS network

```
PE1#show ip vrf interfaces
```

Interface	IP-Address	VRF	Protocol
Se2/0	11.1.1.2	one	up
Lo999	200.1.1.1	two	up

```
PE1#show ip vrf
```

Name	Default RD	Interfaces
one	1:1	Se2/0
three	3.3.3.3:3	
two	1:2	Lo999

VRF RIB

PE1# **show ip route vrf one**

Routing Table: one

Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
E1 - OSPF external type 1, E2 - OSPF external type 2
I - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
ia - IS-IS inter area, * - candidate default, U - per-user static route
o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

```
11.0.0.0/8 is variably subnetted, 6 subnets, 2 masks
B    11.1.2.0/24    [200/0] via 10.100.1.4, 2d18h
B    11.1.3.0/24    [200/0] via 10.100.1.6, 2d18h
C    11.1.1.0/24 is directly connected, Serial2/0
B    11.100.1.7/32  [200/1] via 10.100.1.6, 2d18h
B    11.100.1.5/32  [200/1] via 10.100.1.4, 2d18h
R    11.100.1.1/32  [120/1] via 11.1.1.1, 00:00:05, Serial2/0
```

VRF routing tables are normal routing tables, but the next hop IP address can be in global routing table

VRF FIB

PE1#show ip cef vrf one

Prefix	Next Hop	Interface
0.0.0.0/0	drop	Null0 (default route handler entry)
0.0.0.0/32	receive	
11.1.1.0/24	attached	Serial2/0
11.1.1.0/32	receive	
11.1.1.2/32	receive	
11.1.1.255/32	receive	
11.1.2.0/24	10.1.1.3	Serial3/0
11.1.3.0/24	10.1.5.6	Ethernet0/0
11.100.1.1/32	11.1.1.1	Serial2/0
11.100.1.5/32	10.1.1.3	Serial3/0
11.100.1.7/32	10.1.5.6	Ethernet0/0

PE1#show ip cef vrf one 11.1.2.0

11.1.2.0/24, version 9, epoch 0, cached adjacency to Serial3/0
0 packets, 0 bytes
tag information set
local tag: VPN-route-head
fast tag rewrite with Se3/0, point2point, tags imposed: {19 22}
via 10.100.1.4, 0 dependencies, recursive
next hop 10.1.1.3, Serial3/0 via 10.100.1.4/32
valid cached adjacency
tag rewrite with Se3/0, point2point, tags imposed: {19 22}

RD

- Makes customer IPv4 prefix unique
- RD is present in the NLRI (MP_REACH_NLRI or MP_UNREACH_NLRI), together with the IPv4 prefix and MPLS label
- RD = 64 bits is added to make VPNv4 prefix unique
- RD comprises Administrator subfield:Assigned number subfield
- Two formats

```
PE1 (config) #ip vrf three
```

```
PE1 (config-vrf) #rd ?
```

ASN:nn or IP-address:nn VPN Route Distinguisher

```
!  
ip vrf three  
rd 3.3.3.3:3  
route-target export 1:3  
route-target import 1:3  
!  
ip vrf two  
rd 1:2  
route-target export 1:2  
route-target import 1:2
```

RD-type 1

RD-type 0

Route-Targets

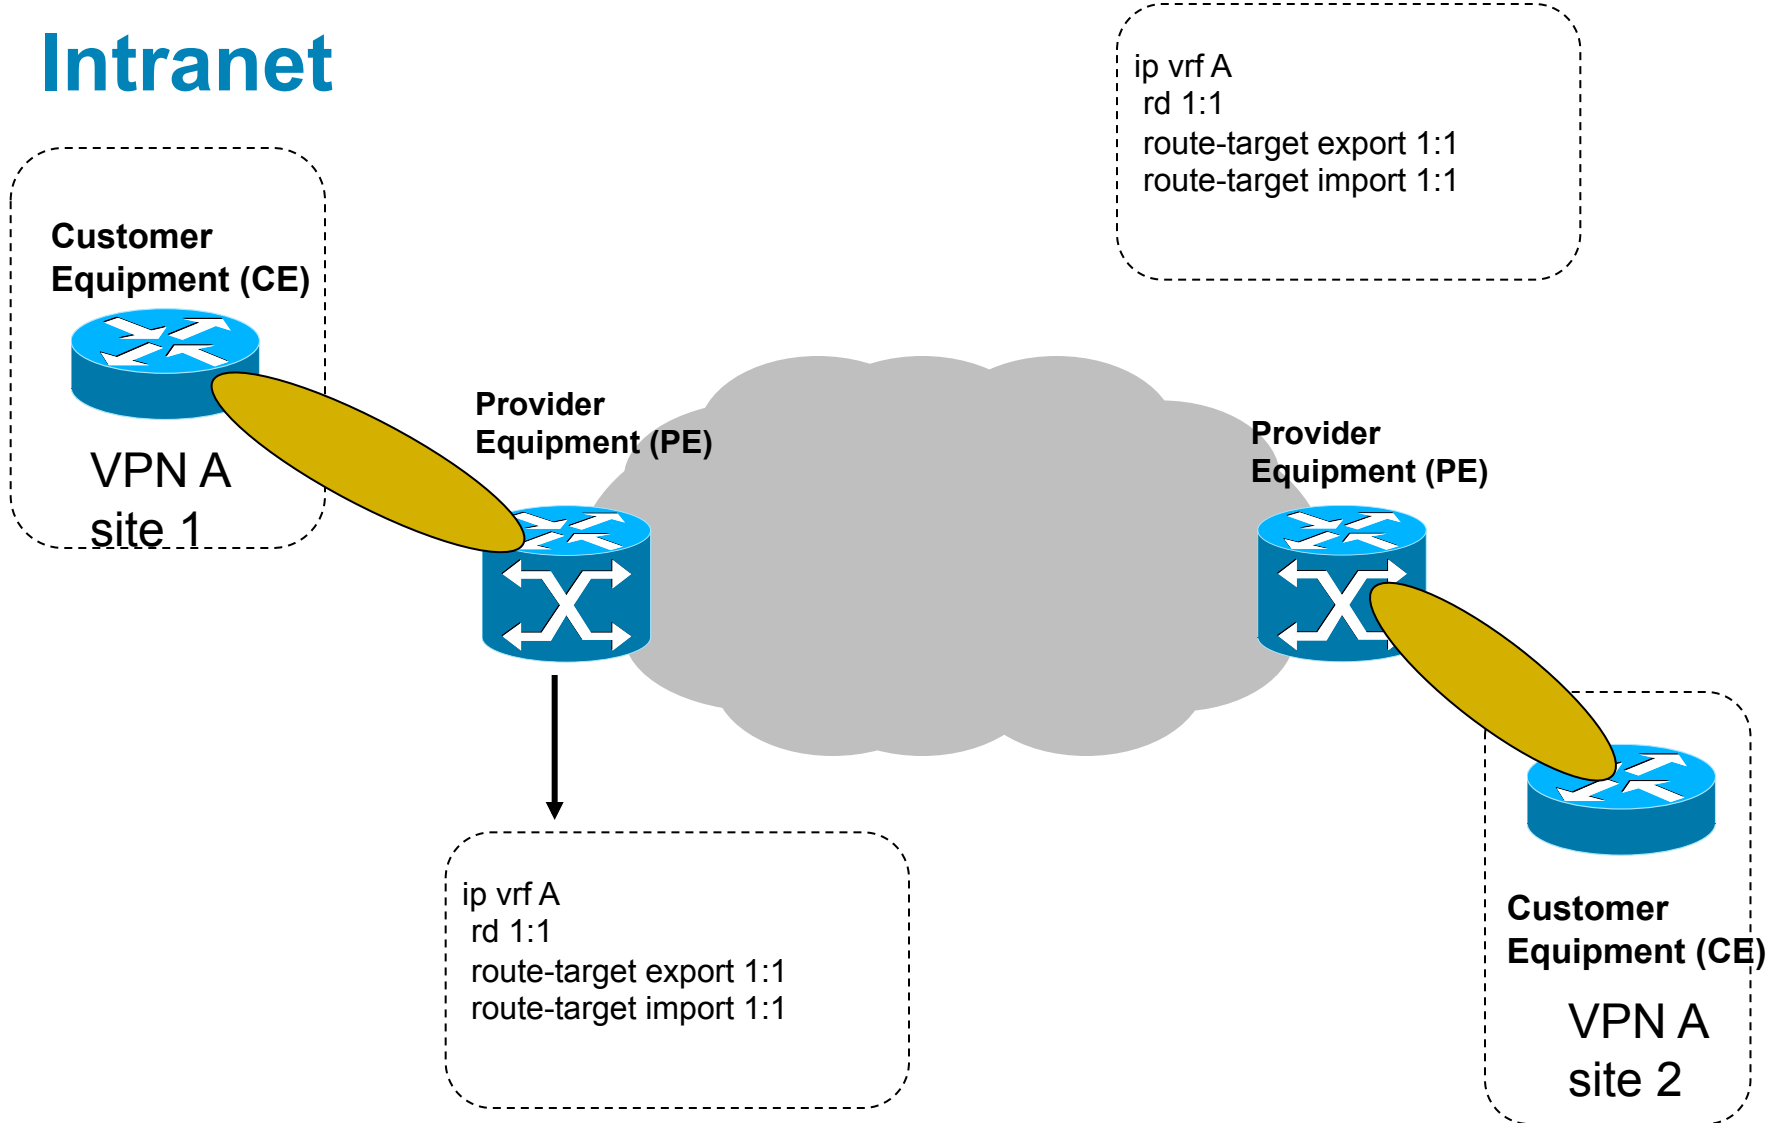
- **Operation**

- Used to control which routes are imported into which VRFs from the remote PE routers and with which Route Targets the vpnv4 routes are exported towards the remote PE routers
- There could be more than one Route Target attached to the vpnv4 route
- For the import into the VRF to be permitted, only one Route Target from the vpnv4 route needs to be matched with the configuration of the imported Route Targets under the *ip vrf* section on the PE router

- **Exporting a Route Target (RT)** means that the exported vpnv4 route will receive an additional BGP extended community (this is the Route Target) as configured under *ip vrf* on the PE router, when the route is redistributed from the VRF routing table into MP-BGP.

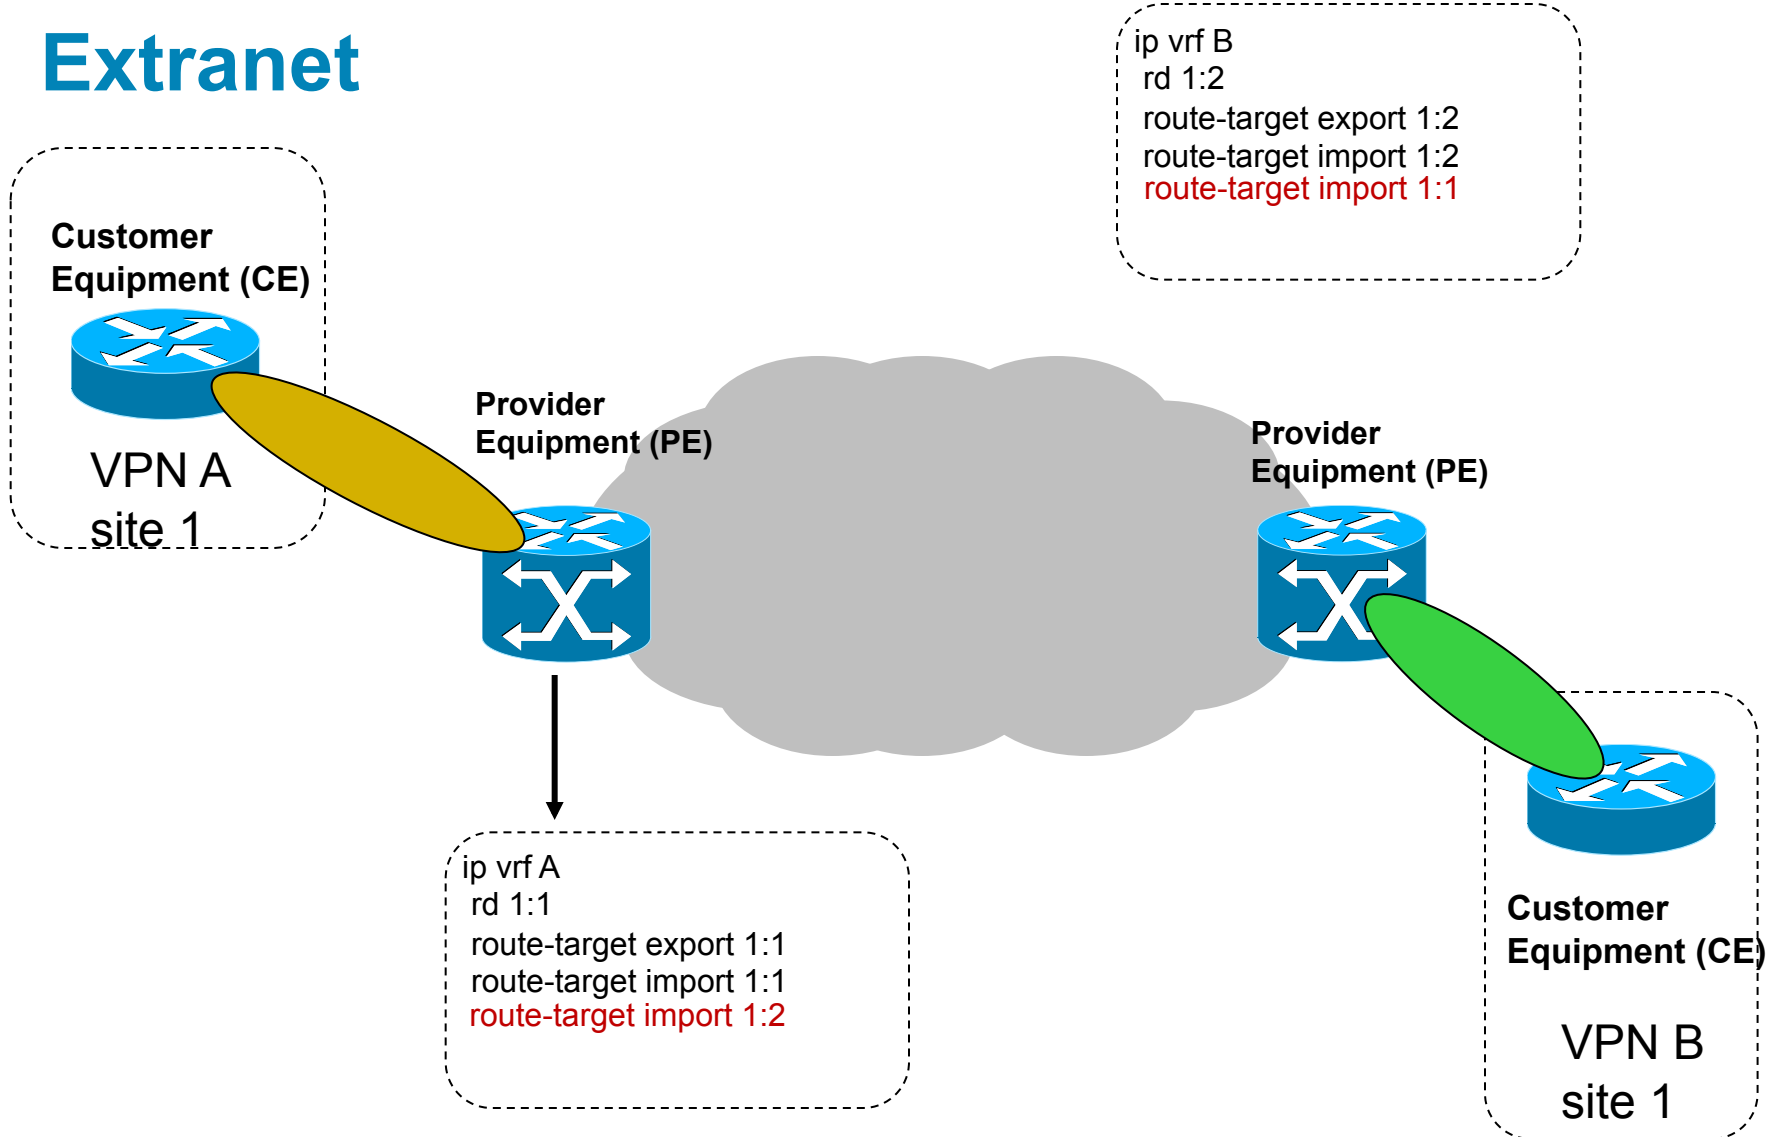
- **Importing a Route Target (RT)** means that the received vpnv4 route from MP-BGP is checked for a matching extended community (this is the route target) with the one in the configuration.

Intranet



Prefixes from VPN A site 1 will be imported into site 2 of VPN A and vice versa

Extranet

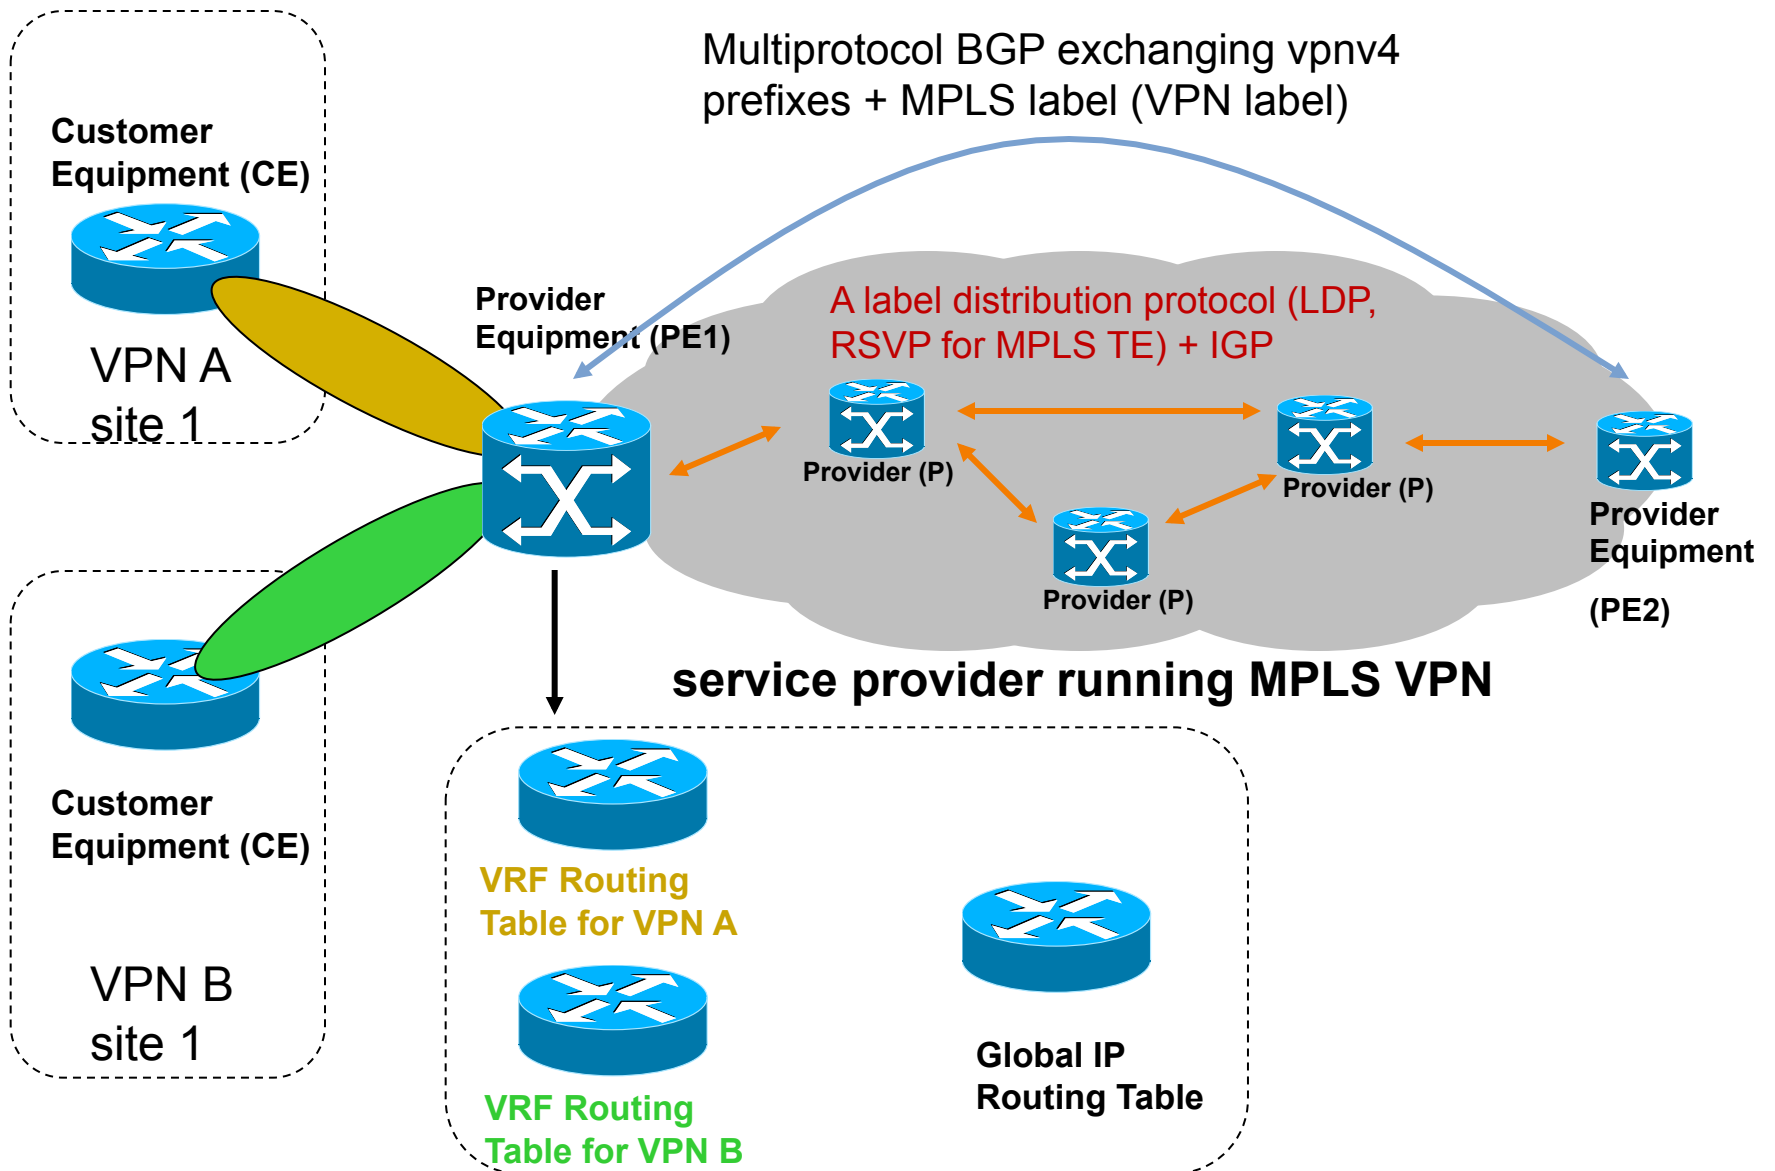


Prefixes from VPN A site 1 will be imported into site 1 of VPN B and vice versa

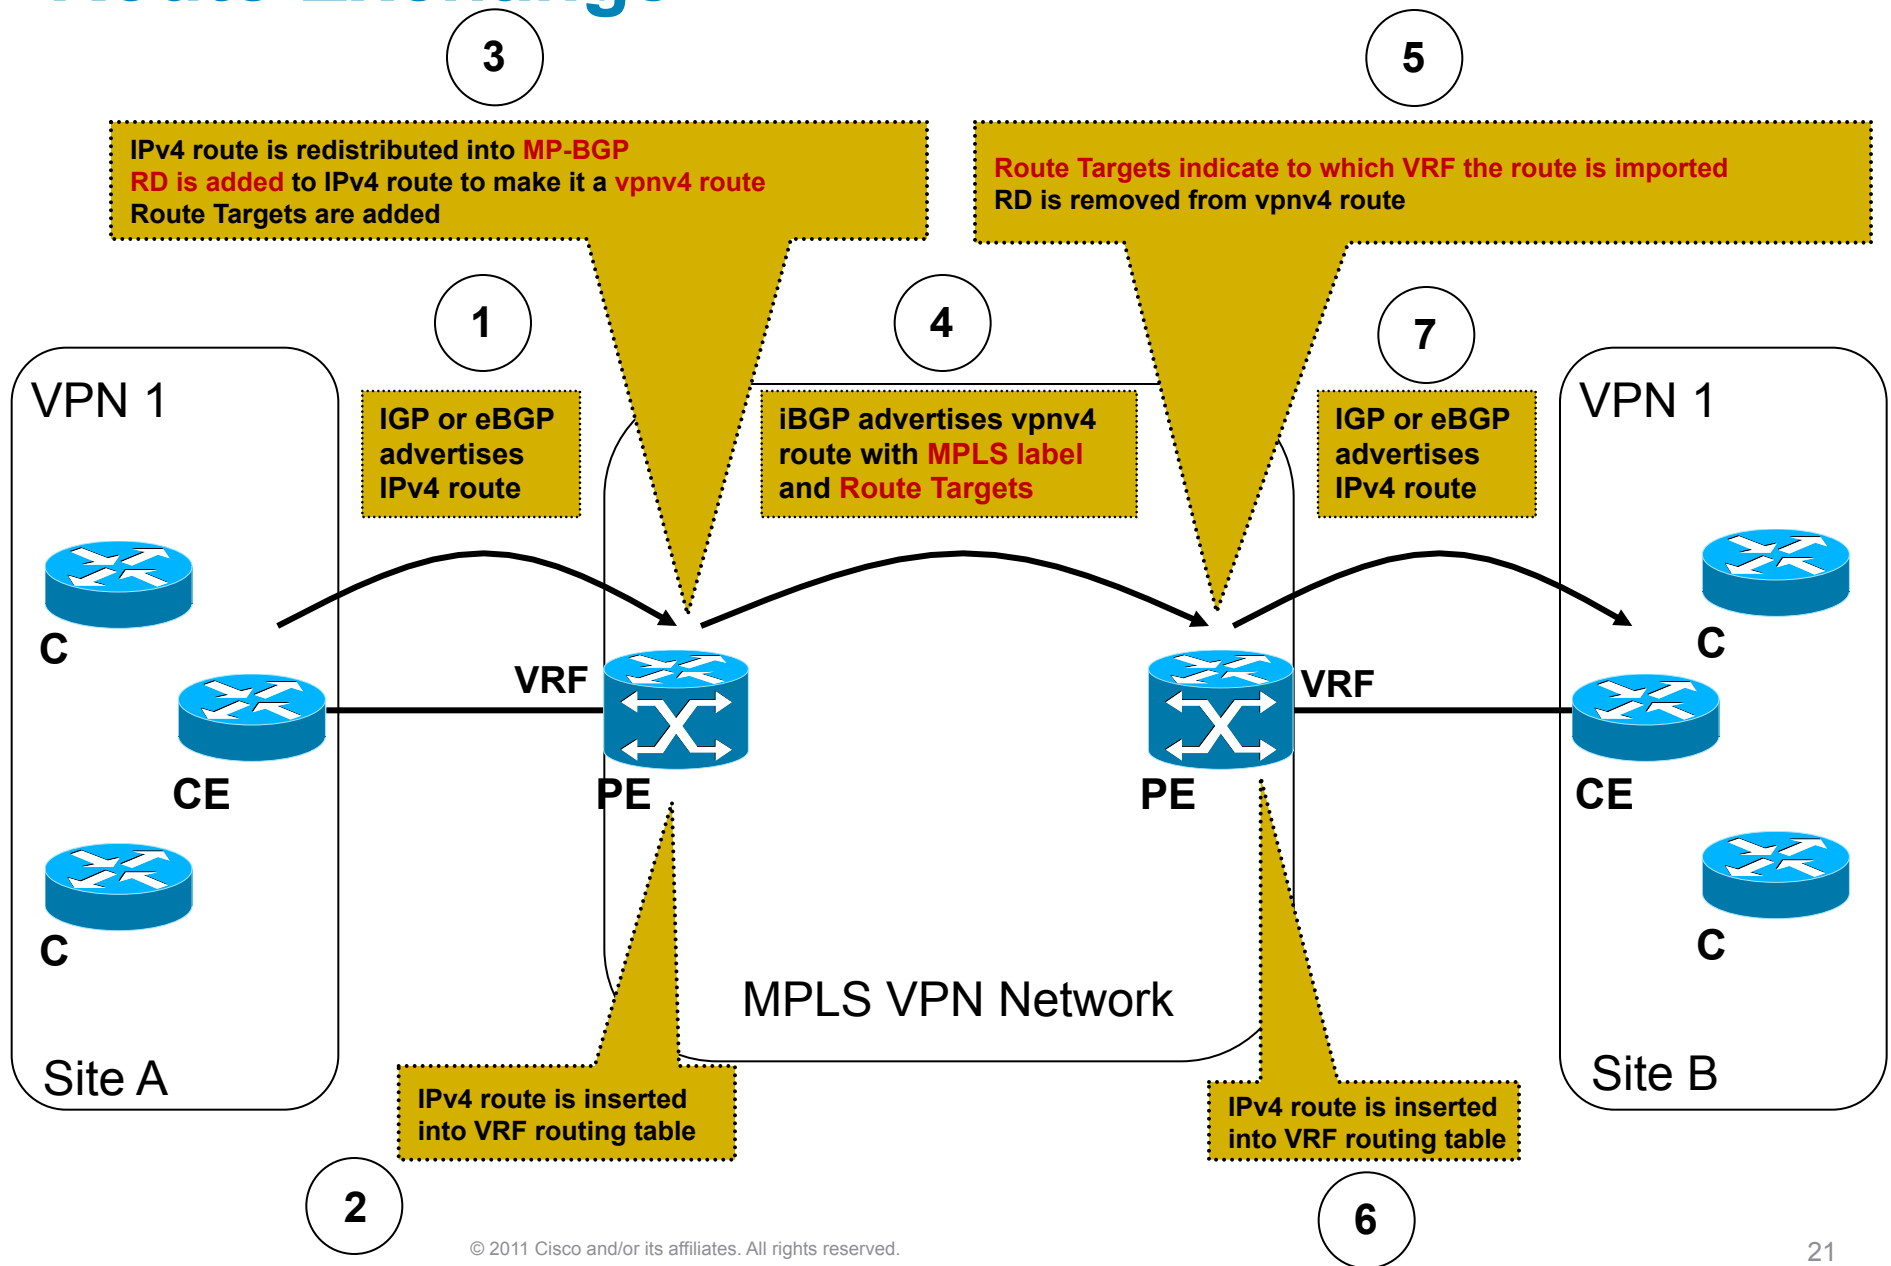
Role of BGP

- iBGP carries:
 - the vpnv4 prefix
 - $\text{vpn4 prefix} = \text{RD} + \text{IPv4 prefix}$
 - Route Target (RT)
 - Any other community and BGP attribute
 - The MPLS label
- Address-family (AF) vpnv4 is used
- Label is automatically advertised in AF vpnv4

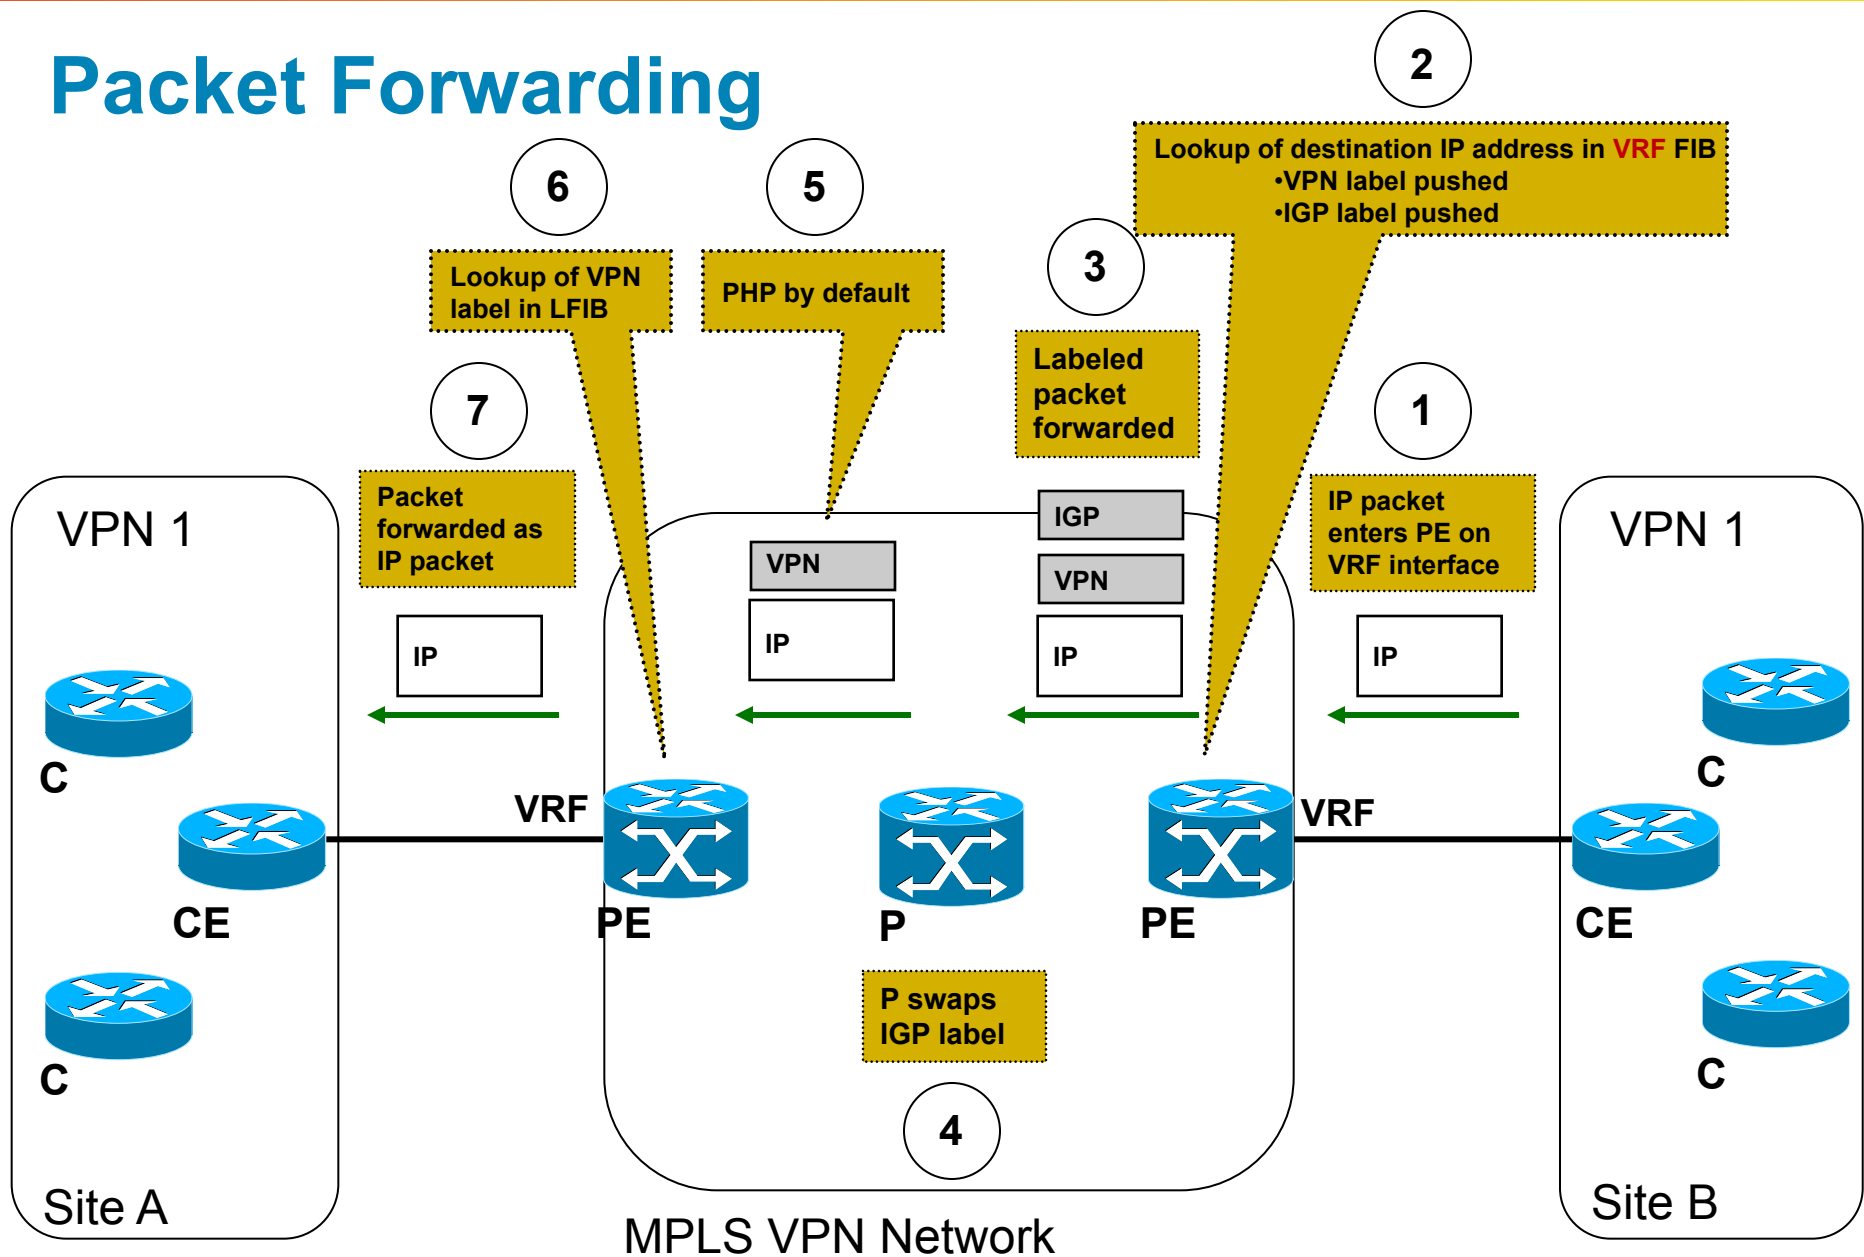
Context




Route Exchange



Packet Forwarding

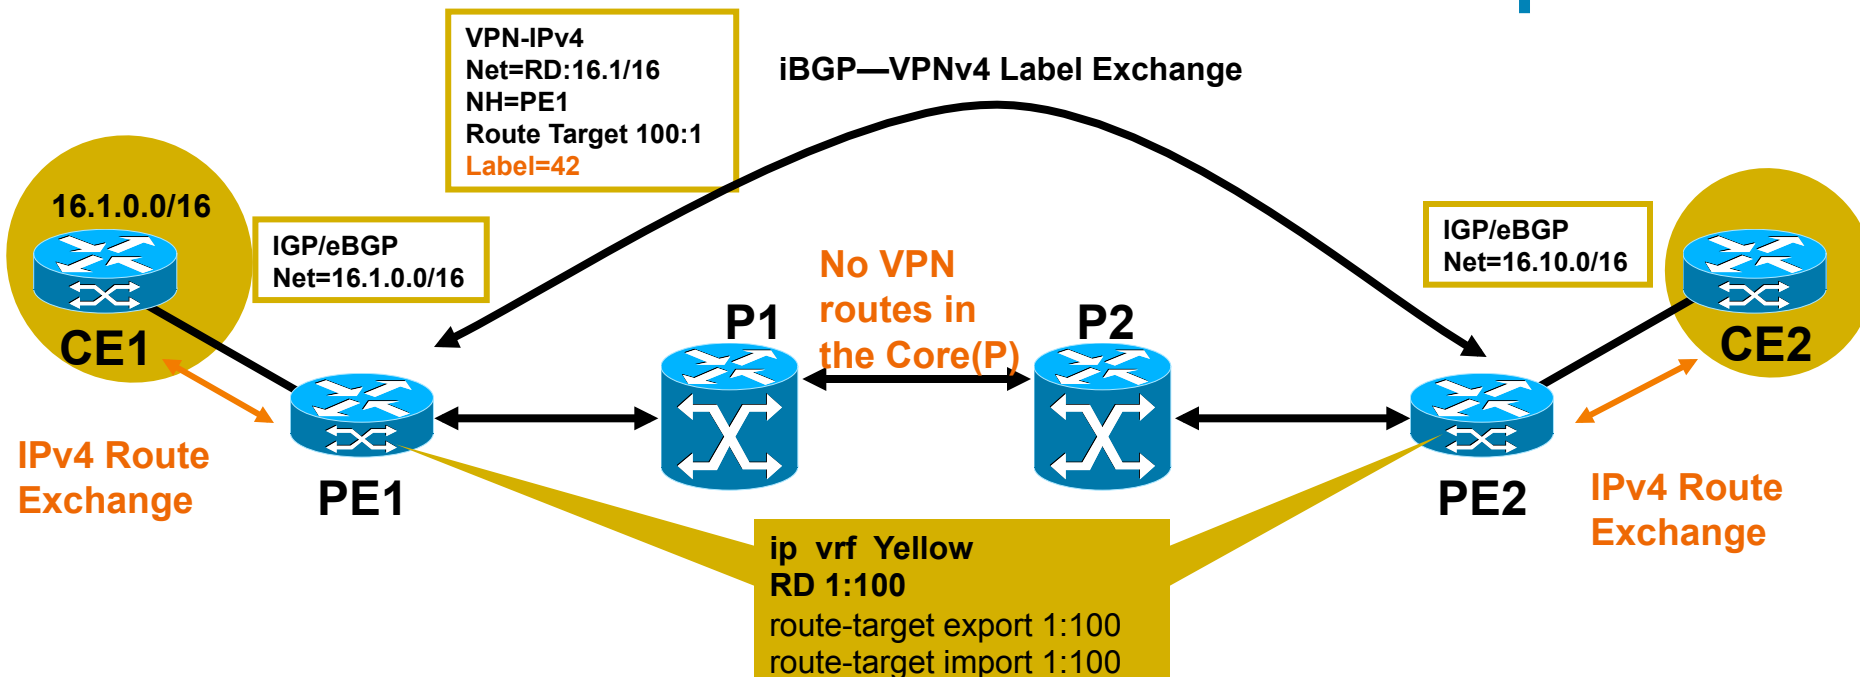


MPLS L3 VPN Control Plane Basics



1. VPN service is enabled on PEs (VRFs are created and applied to VPN site interface)
 2. VPN site's CE1 connects to a VRF enabled interface on a PE1
 3. VPN site routing by CE1 is distributed to MP-iBGP on PE1
 4. PE1 allocates VPN label for each prefix, sets itself as a next hop and relays VPN site routes to PE3
 5. PE3 distributes CE1's routes to CE2
- (Similar happens from CE2 side...)

How Control Plane Information Is Separated

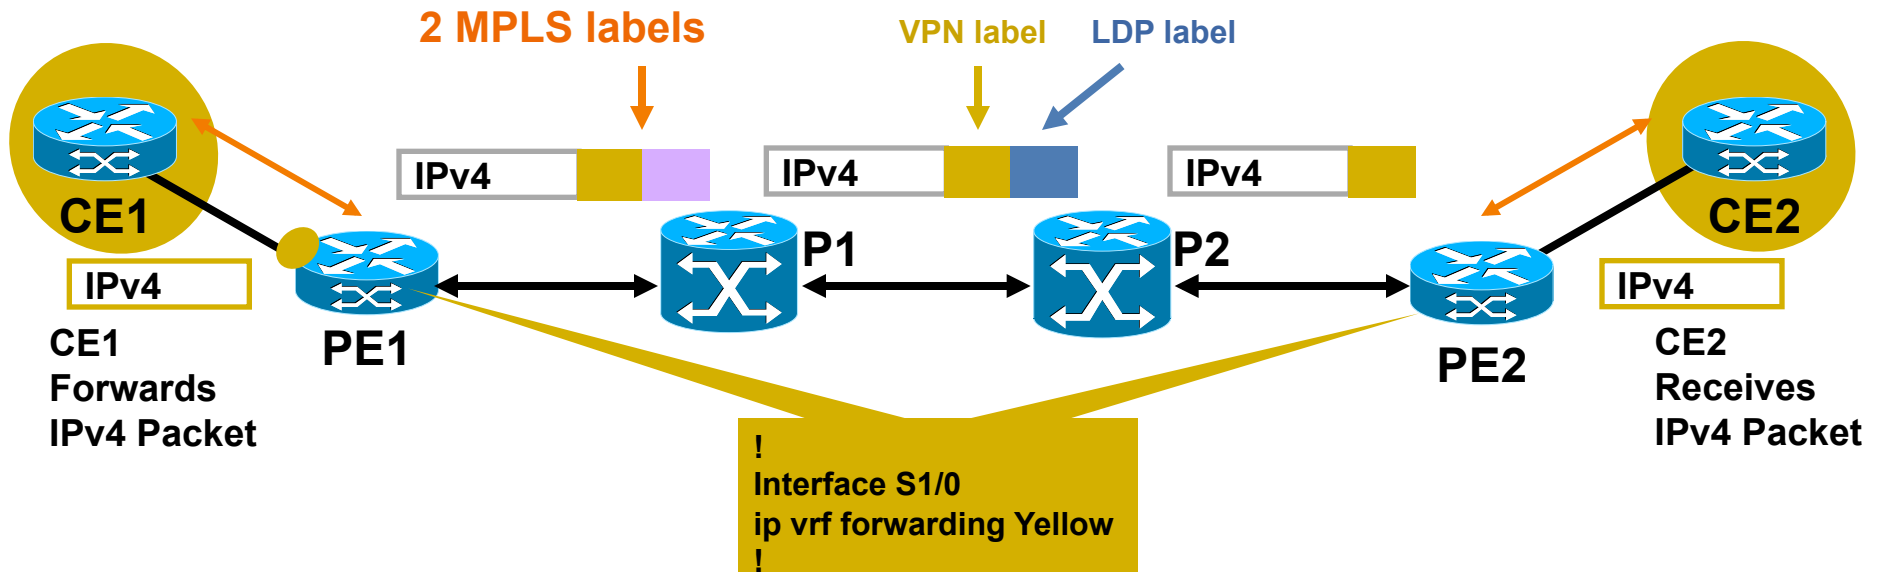


MPLS VPN Control Plane Components:

- **Route Distinguisher:** 8 byte field—unique value assigned by a provider to each VPN to make a route unique so customers don't see each other's routes
- **VPNv4 address:** RD+VPN IP prefix;
- **Route Target:** RT-8bytes field, unique value assigned by a provider to define the import/export rules for the routes from/to each VPN
- **MP-BGP:** facilitates the advertisement of VPNv4* prefixes + labels between MP-BGP peers
- **Virtual Routing Forwarding Instance (VRF):** contains VPN site routes
- **Global Table:** Contains core routes, Internet or routes to other services

How Does It Work?

How Data Plane Is Separated



1. PE1 imposes pre allocated label for the prefix
2. Core facing interface allocates IGP label
3. Core swap IGP labels
4. PE2 strips off VPN label and forwards the packet to CE2 as an IP packet

Config on PE Router

Definition of VRF

Assigning CE-facing interface to VRF

PE-CE routing protocol

Configuring BGP vpnv4 peering
(to all other PEs or RRs)

Configuring VRF BGP (redistribution)

```
ip vrf one
rd 1:1
route-target export 1:1
route-target import 1:1
!
interface FastEthernet2/1
ip vrf forwarding one
ip address 99.1.1.2 255.255.255.0
!
router ospf 100 vrf one
log-adjacency-changes
redistribute bgp 1 metric 10 subnets
network 99.1.1.0 0.0.0.255 area 0
!
router bgp 1
bgp log-neighbor-changes
neighbor 11.100.100.4 remote-as 1
neighbor 11.100.100.4 update-source Loopback0
!
address-family ipv4
no synchronization
neighbor 11.100.100.4 activate
neighbor 11.100.100.4 send-community both
exit-address-family
!
address-family vpnv4
neighbor 11.100.100.4 activate
neighbor 11.100.100.4 send-community both
exit-address-family
!
address-family ipv4 vrf one
redistribute connected
redistribute ospf 100 vrf one
exit-address-family
```

Verifying VPNv4 Prefixes in BGP

PE1#**show ip bgp vpnv4 all**

BGP table version is 15, local router ID is 10.100.1.2

Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
r RIB-failure, S Stale

Origin codes: i - IGP, e - EGP, ? - incomplete

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 1:1 (default for vrf one)					
*> 11.1.1.0/24	0.0.0.0	0	32768	?	
*>i11.1.2.0/24	10.100.1.4	0	100	0	?
*>i11.1.3.0/24	10.100.1.6	0	100	0	?
*> 11.100.1.1/32	11.1.1.1	1	32768	?	
*>i11.100.1.5/32	10.100.1.4	1	100	0	?
*>i11.100.1.7/32	10.100.1.6	1	100	0	?
Route Distinguisher: 1:2 (default for vrf two)					
*>i14.1.1.1/32	10.100.1.4	0	100	0	?

PE1#**debug ip bgp vpnv4 unicast updates**

BGP updates debugging is on for address family: VPNv4 Unicast

BGP(2): 10.100.1.4 rcvd UPDATE w/ attr: nexthop 10.100.1.4, origin ?, localpref 100, metric 0, extended community RT:
1:2

BGP(2): 10.100.1.4 **rcvd 1:2:14.1.1.1/32**

MPLS Aware ICMP

```
PE1#trace vrf one 11.100.1.5
```

Type escape sequence to abort.

Tracing the route to 11.100.1.5

```
 1 10.1.1.3 [MPLS: Labels 19/23 Exp 0] 32 msec 60 msec 40 msec
 2 11.1.2.4 [MPLS: Label 23 Exp 0] 40 msec 20 msec 20 msec
 3 11.1.2.5 60 msec * 64 msec
```

ICMP in IOS can carry the label stack when generating ICMP reply messages

PE-CE Routing Protocols

- Connected
- Static
- RIPv2
- OSPF
- EIGRP
- eBGP

VRF Access

- Cisco IOS commands were made VRF aware in order to be able to communicate with the CE devices or IP addresses on the PE router in the VRF context

```
london# ping vrf cust-one 10.10.100.1
```

```
london# traceroute vrf cust-one 10.10.100.1
```

```
london# telnet 10.10.100.1 /vrf cust-one
```

MPLS VPN Troubleshooting

- Basic Checks:
 - ping and traceroute CE-to-CE
 - Ping local PE-CE
 - Ping from PE to remote PE-CE
 - Ping and traceroute PE-to-PE in global RIB (between PE loopback)
- Ping check connectivity
- Traceroute can tell us if LSP is broken (look for missing label)
- If there is a failure:
 - check routing tables (global and VRF)
 - Check forwarding vector
 - Check CEF tables (global and VRF)
 - Check forwarding vector and labels
 - Check BGP vpnv4 table
 - Check next-hop and labels

MPLS VPN Troubleshooting: Example

CE1#ping 11.100.1.5

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 11.100.1.5, timeout is 2 seconds:

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 60/60/64 ms

CE1#trace 11.100.1.5

Type escape sequence to abort.

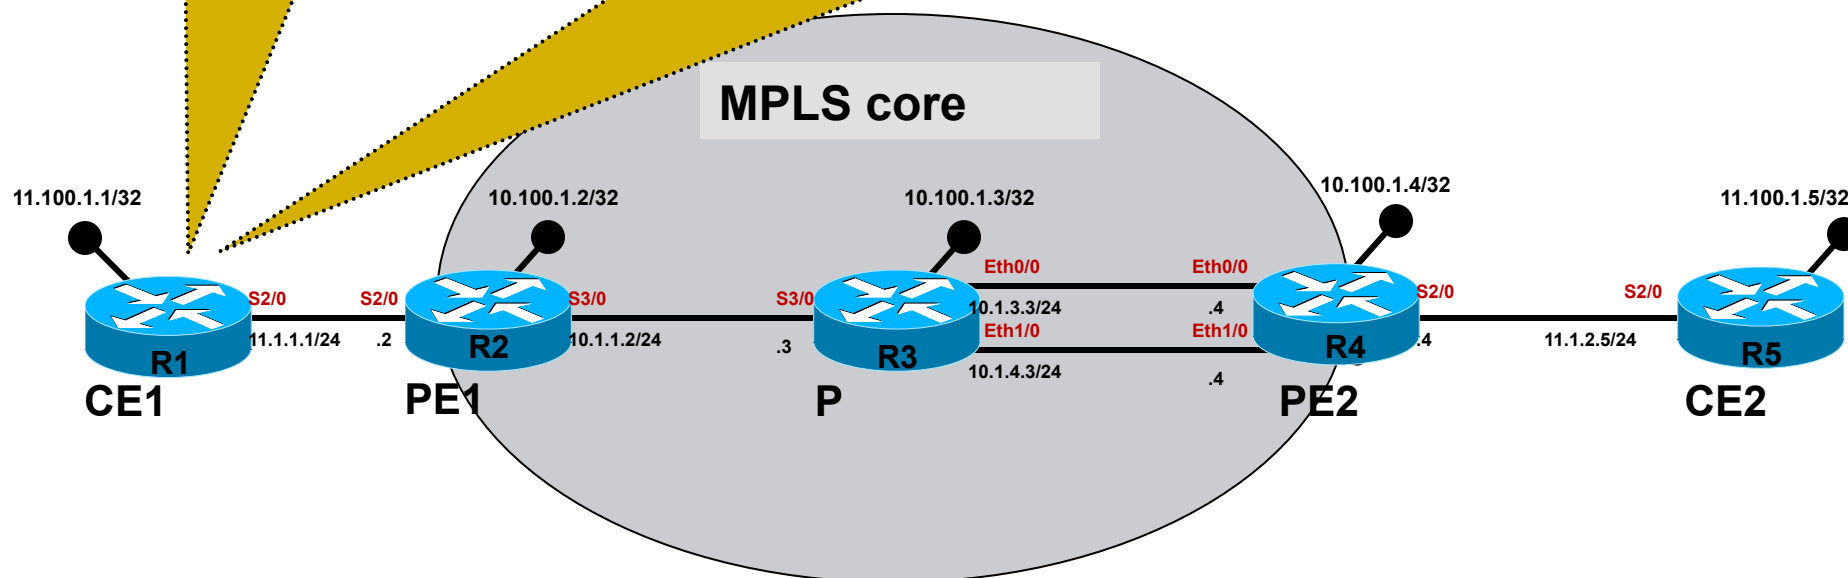
Tracing the route to 11.100.1.5

1 11.1.1.2 20 msec 20 msec 20 msec

2 10.1.1.3 [MPLS: Labels 19/23 Exp 0] 60 msec 60 msec 60 msec

3 11.1.2.4 [MPLS: Label 23 Exp 0] 40 msec 40 msec 40 msec

4 11.1.2.5 48 msec * 64 msec



MPLS VPN Troubleshooting: Example

PE1#ping vrf one 11.100.1.5

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 11.100.1.5, timeout is 2 seconds:

!!!!


Success rate is 100 percent (5/5), round-trip min/avg/max = 40/52/72 ms

PE1#trace vrf one 11.100.1.5

Type escape sequence to abort.

Tracing the route to 11.100.1.5

```
1 10.1.1.3 [MPLS: Labels 19/23 Exp 0] 32 msec 60 msec 40 msec
2 11.1.2.4 [MPLS: Label 23 Exp 0] 40 msec 20 msec 20 msec
3 11.1.2.5 60 msec * 64 msec
```




MPLS VPN Troubleshooting: Example

```
PE1#show ip cef vrf one 11.100.1.5
11.100.1.5/32, version 10, epoch 0, cached adjacency to
Serial3/0 0 packets, 0 bytes
tag information set
  local tag: VPN-route-head
  fast tag rewrite with Se3/0, point2point, tags imposed: {19 23}
via 10.100.1.4, 0 dependencies, recursive
next hop 10.1.1.3, Serial3/0 via 10.100.1.4/32
valid cached adjacency
tag rewrite with Se3/0, point2point, tags imposed: {19 23}
```

```
PE1#show ip cef 10.100.1.4
10.100.1.4/32, version 20, epoch 0, cached adjacency to
Serial3/0 0 packets, 0 bytes
tag information set, shared
  local tag: 23
  fast tag rewrite with Se3/0, point2point, tags imposed: {19}
via 10.1.1.3, Serial3/0, 2 dependencies
next hop 10.1.1.3, Serial3/0
valid cached adjacency
tag rewrite with Se3/0, point2point, tags imposed: {19}
```


```
PE1#show ip bgp vpnv4 vrf one 11.100.1.5
BGP routing table entry for 1:1:11.100.1.5/32, version 12
Paths: (1 available, best #1, table one)
Not advertised to any peer
Local
  10.100.1.4 (metric 75) from 10.100.1.4 (10.100.1.4)
Origin incomplete, metric 0, localpref 100, valid, internal, best
Extended Community: RT:1:1
mpls labels in/out nolabel/23
```



MPLS VPN Troubleshooting: Example

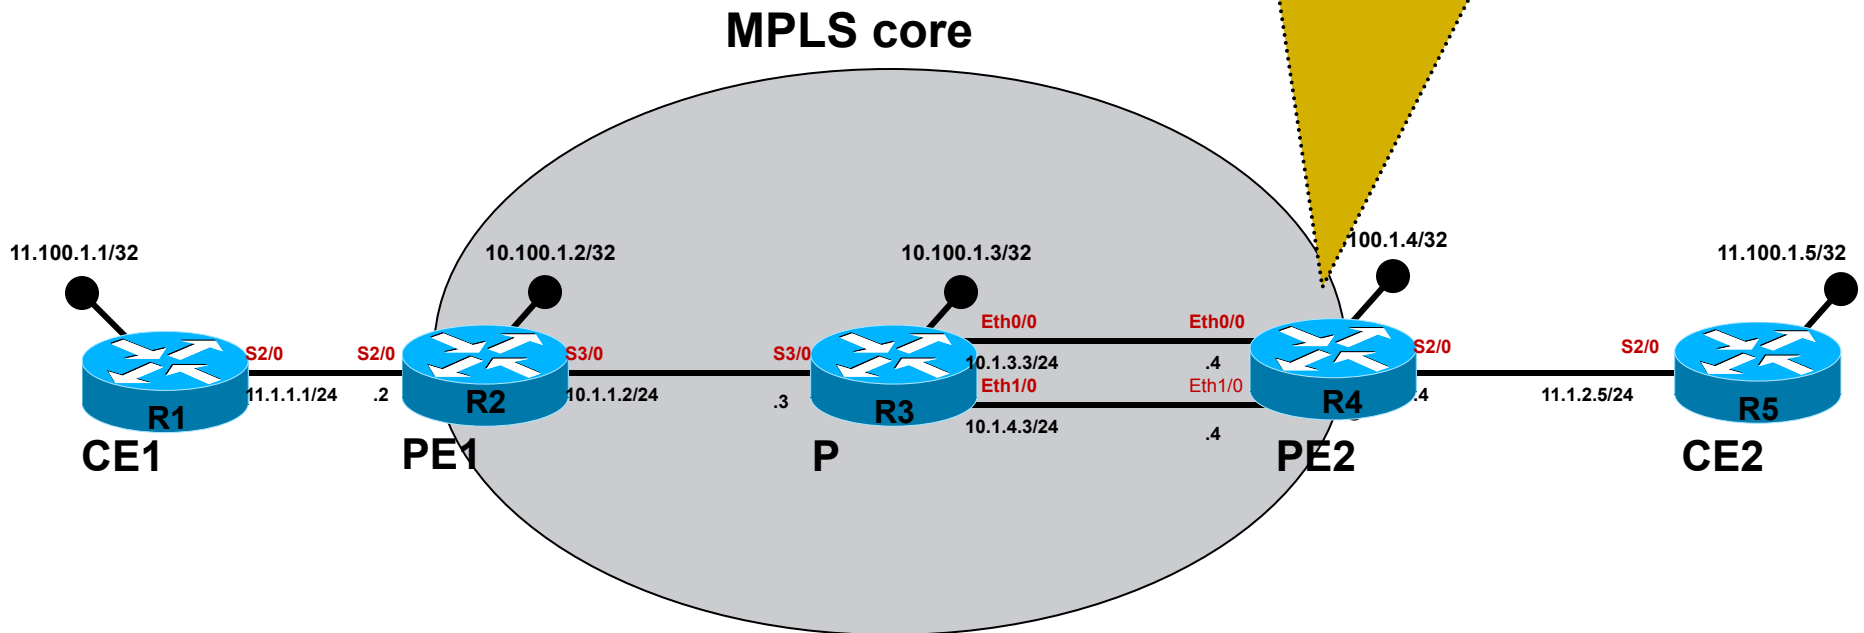
```
P#show mpls forwarding-table 10.100.1.4
```

Local tag	Outgoing tag or VC	Prefix or Tunnel Id	Bytes switched	Outgoing interface	Next Hop
19	Pop tag	10.100.1.4/32	1220	Et1/0	10.1.4.4
	Pop tag	10.100.1.4/32	1051348	Et0/0	10.1.3.4



MPLS VPN Troubleshooting: Example

```
PE2#show mpls forwarding-table labels 23
Local  Outgoing  Prefix          Bytes Label  Outgoing  Next Hop
Label  Label or VC  or Tunnel Id    Switched     interface
23     No Label    11.100.1.5/32[V] 0             Se2/0      point2point
```



MPLS VPN Troubleshooting: Example

Debugging Control plane

```
PE1#debug ip bgp vpnv4 unicast updates
```

BGP updates debugging is on for address family: VPNv4 Unicast

```
PE1#show debug
```

IP routing:

BGP updates debugging is on for address family: VPNv4 Unicast

```
PE2#debug ip bgp vpnv4 unicast updates
```

BGP updates debugging is on for address family: VPNv4 Unicast

```
PE2#show debug
```

IP routing:

BGP updates debugging is on for address family: VPNv4 Unicast

```
CE1#debug ip routing
```

IP routing debugging is on

```
CE1#show debug
```

IP routing:

IP routing debugging is on

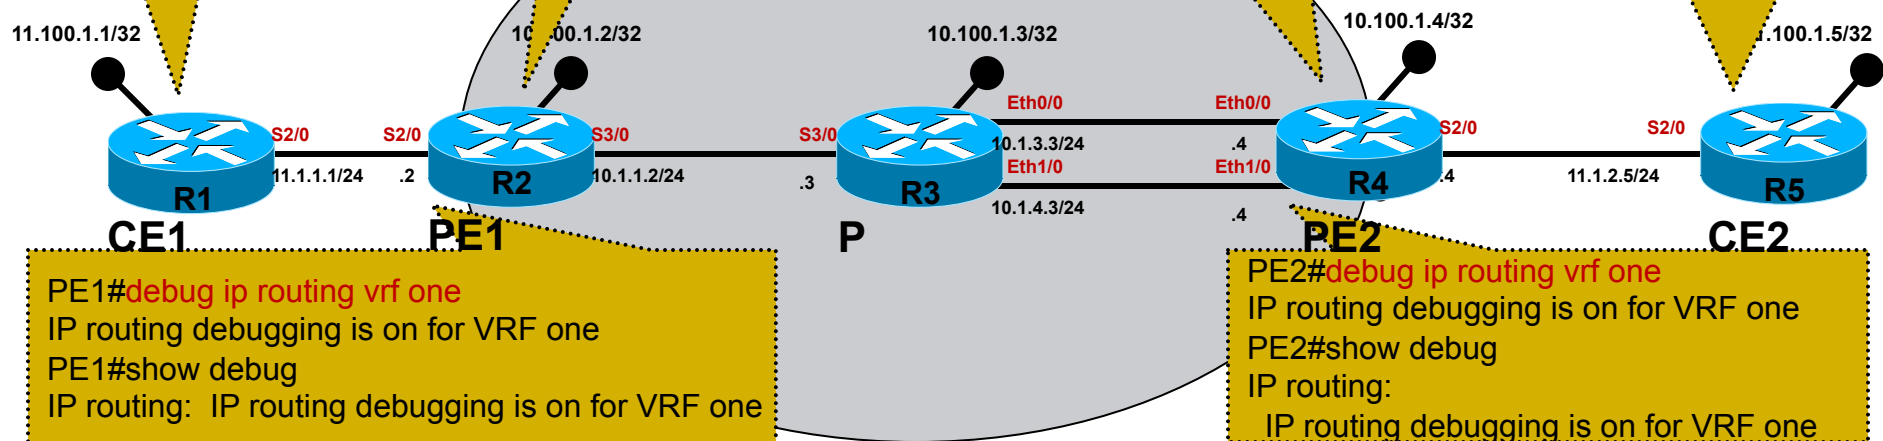
```
CE2#debug ip routing
```

IP routing debugging is on

```
CE2#show debug
```

IP routing:

IP routing debugging is on

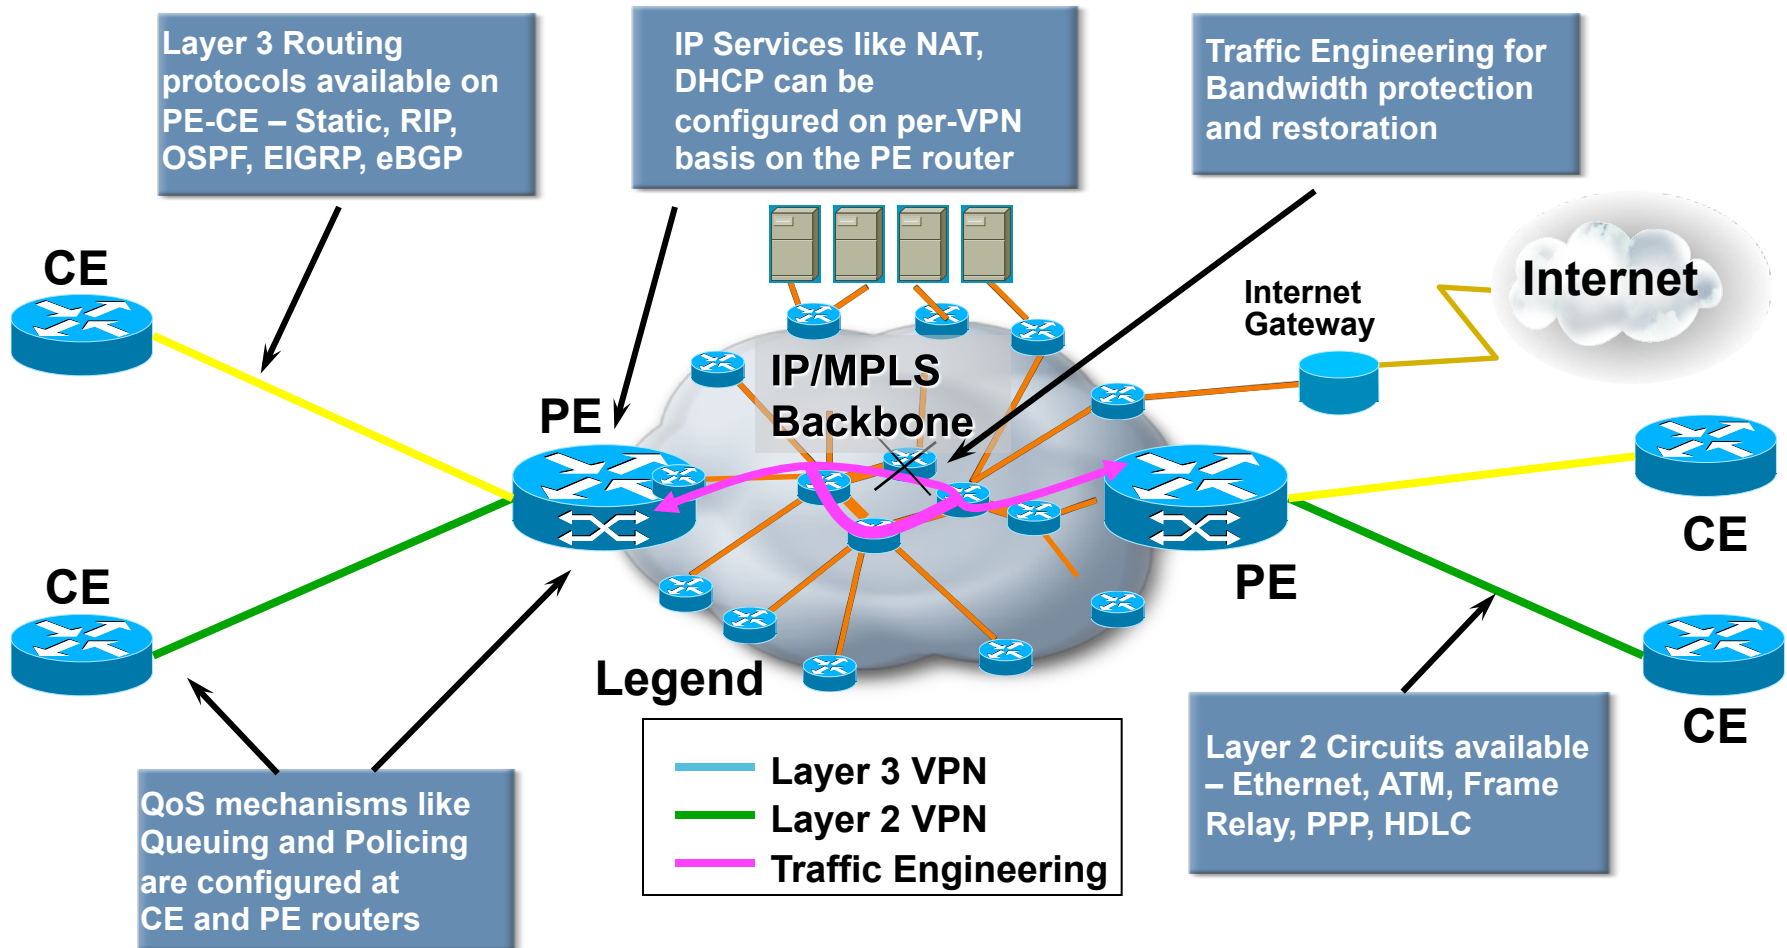




Summary



The Full Service Network: Integrated MPLS Technologies




Layer 3 VPNs & Layer 2 VPNs, Traffic Engineering + QoS + IP Services

Ask The Experts Event (with Nagendra Kumar)

If you have additional questions, you can ask them to Nagendra here:

<https://supportforums.cisco.com/community/netpro/ask-the-expert>

He will be answering from August 17th to August 26th.



First CSC Expert Series Webcast in Spanish

Topic: Firewall Service Module: Architecture and Operation

Tuesday, August 30th, at

7:00 a.m. Pacific (UTC -7)

9:00 a.m Mexico city (UTC -7)

4:00 p.m Madrid (UTC +2)

Join Security CCIE and Certified Ethical Hacker from EC-Council

Ivan Martiñón from **HTTS** group in **Latin America**

He will talk about the architecture, operation and configuration of the Firewall Service Modules, as well as Firewall technologies in general.

During this interactive session you will be able ask all your questions related to this topic.

Register for this live Webcast at

http://bitly.com/webcast_registration



Next CSC Expert Series Webcast

Topic: Things I Can Do to Protect My Network from Getting Hacked

Tuesday, September 13th, at
8:00 a.m. Pacific Time
6:00 p.m. CEST Brussels (UTC +2),
11:00 a.m. EDT New York (UTC -4).



Join double CCIE, Technical Leader

Jazib Frahim from **RTP**.

He will provide reasons why enterprise network segments get compromised despite their state-of-the-art network security technologies and products that are deployed.

During this interactive session you will be able ask all your questions related to this topic.

Register for this live Webcast at

www.CiscoLive.com/ATE

Did you know that we have communities in other languages?

If you speak **Polish, Japanese, or Spanish**, we invite you to ask your questions and collaborate in your language.

- **Spanish** → <https://supportforums.cisco.com/community/spanish>
- **Polish** → <https://supportforums.cisco.com/community/etc/netpro-polska>
- **Japanese** → <https://supportforums.cisco.com/community/csc-japan>

We're also running a pilot for **Russian and Portuguese**. You can register at the following links

- **Russian:**
<https://www.ciscofeedback.vovici.com/se.ashx?s=6A5348A712220E19>
- **Portuguese:**
<https://www.ciscofeedback.vovici.com/se.ashx?s=6A5348A77EE5C0B7>

